

# Artificial Intelligence Frameworks

## Table of Contents

1. Artificial Intelligence
2. Data Science
3. Machine Learning
4. Deep Learning
5. Natural Language Processing (NLP)
6. Large Language Models (LLMs)
7. Generative AI
8. AI Agents / Agentic AI
9. Computer Vision (CV)
10. Evaluation
11. Optimization
12. Deployment
13. MLOps
14. Hardware Acceleration

## 1 — ARTIFICIAL INTELLIGENCE (General AI Development)

Rank	Framework/Tool	Description
1	Python	Primary programming language for all AI development.
2	NumPy	Fast numerical computation, arrays, matrices.
3	Pandas	Data manipulation & analysis with DataFrames.
4	SciPy	Scientific computing, optimization, statistics.
5	Scikit-learn	Classical ML algorithms & preprocessing utilities.
6	Jupyter Notebook	Interactive coding environment for AI experiments.
7	ONNX / ONNX Runtime	Standard for model interchange & optimized inference.

8	Ray	Distributed computing for ML & AI workloads.
9	Docker	Containerization for deploying AI models.
10	FastAPI	High-performance API framework to serve AI models.

---

## 2 — DATA SCIENCE

Rank	Framework/Tool	Description
1	Pandas	Data manipulation, wrangling, cleaning.
2	NumPy	Numerical computing backbone of Python DS stack.
3	Matplotlib	Base data visualization library.
4	Seaborn	Statistical visualizations built on Matplotlib.
5	Plotly	Interactive charts, dashboards.
6	Polars	Next-gen ultra-fast DataFrame library.
7	Apache Spark (PySpark)	Distributed processing for big data.
8	Dask	Parallel computing on local or cluster machines.
9	Airflow	Workflow orchestration for data pipelines.
10	SQLAlchemy	ORM & database query interface.

---

## 3 — MACHINE LEARNING (Classical ML)

Rank	Framework/Tool	Description
1	Scikit-learn	Full suite of ML algorithms & preprocessing.
2	XGBoost	High-performance gradient boosting.
3	LightGBM	Fast GPU-optimized gradient boosting.
4	CatBoost	Best boosting library for categorical data.
5	TensorFlow	ML + DL framework for scalable models.
6	PyTorch	Most popular DL framework (also ML support).
7	H2O.ai AutoML	Automated ML pipeline generation.
8	MLflow	Experiment tracking & model registry.

9	Optuna	Hyperparameter optimization library.
10	StatsModels	Statistical & econometric modeling.

---

## 4 — DEEP LEARNING

Rank	Framework/Tool	Description
1	PyTorch	Most-used DL library for research & production.
2	TensorFlow	Enterprise-grade DL framework.
3	Keras	High-level API for building DL models quickly.
4	JAX	High-performance ML with auto-differentiation.
5	HuggingFace Transformers	Pretrained DL models for NLP, CV, audio.
6	FastAI	Simplified high-level DL training.
7	DeepSpeed	Microsoft's large-model training framework.
8	Horovod	Distributed training across clusters.
9	OpenVINO	Optimized inference for Intel hardware.
10	TensorRT	NVIDIA GPU inference acceleration.

---

## 5 — NATURAL LANGUAGE PROCESSING (NLP)

Rank	Framework/Tool	Description
1	HuggingFace Transformers	SOTA models for NLP tasks.
2	spaCy	Industrial NLP for production apps.
3	NLTK	Classic NLP processing toolkit.
4	Gensim	Topic modeling & word embeddings.
5	SentenceTransformers	Embedding generation for semantic tasks.
6	fastText	Facebook's fast text classification.
7	Stanford CoreNLP	Java-based advanced NLP toolkit.
8	AllenNLP	Research-oriented NLP framework.
9	Flair	NLP classification & NER framework.

10	OpenAI GPT API	Advanced NLP via LLM API access.
----	----------------	----------------------------------

---

## 6 — LARGE LANGUAGE MODELS (LLMs)

Rank	Framework/Tool	Description
1	HuggingFace Transformers	Core framework for LLM training & inference.
2	OpenAI API	Access to GPT models for LLM applications.
3	LangChain	LLM orchestration & agent framework.
4	LlamaIndex	Document indexing for LLM apps.
5	vLLM	Fastest inference engine for LLMs.
6	TGI (Text Generation Inference)	Production LLM inference server.
7	DeepSpeed	Efficient large model training.
8	PEFT	Parameter-efficient fine-tuning (LoRA/QLoRA).
9	bitsandbytes	Quantization for model size reduction.
10	FlashAttention	Optimized attention computation for LLMs.

---

## 7 — GENERATIVE AI (Image/Video/Audio/Text)

Rank	Framework/Tool	Description
1	Diffusers (HF)	Stable Diffusion pipelines for gen AI.
2	Automatic1111	UI interface for image generation.
3	ComfyUI	Node-based generative AI system.
4	ControlNet	Condition-based image generation.
5	Whisper	SOTA speech-to-text.
6	Tortoise TTS	High-quality text-to-speech.
7	RVC / So-VITS	Voice cloning frameworks.
8	RunwayML	Video generation & editing AI tools.
9	GANs (PyTorch/TensorFlow)	Generative Adversarial Networks.
10	DreamBooth	Personalized fine-tuning for image generation.

---

## 8 — AI AGENTS / AGENTIC AI

Rank	Framework/Tool	Description
1	LangChain Agents	Agent framework with tools/actions.
2	LlamaIndex Agents	Document+tool based agent framework.
3	OpenAI Swarm	Multi-agent orchestration by OpenAI.
4	Microsoft Autogen	Multi-agent automation system.
5	CrewAI	Role-based AI agent coordination.
6	Haystack Agents	Production-ready agent pipelines.
7	AgentGPT	Browser-based automatic agents.
8	SuperAGI	Open-source autonomous agent framework.
9	LiteLLM	Multi-LLM orchestration layer.
10	ReAct / Toolformer	Architectures for reasoning-with-tools.

---

## 9 — COMPUTER VISION (CV)

Rank	Framework/Tool	Description
1	OpenCV	Core CV operations, image processing.
2	torchvision	PyTorch CV toolkit & pretrained models.
3	TensorFlow Vision	TF-based CV components.
4	YOLOv5	High-accuracy object detection.
5	YOLOv8 / YOLO-NAS	Latest ultrafast detection models.
6	Detectron2	Facebook's advanced CV framework.
7	MMDetection	OpenMMLab's modular CV library.
8	Mediapipe	Hand, face, body landmark detection.
9	Albumentations	Best image augmentation library.
10	Open3D	3D vision & point-cloud processing.

---

## 10 — EVALUATION – Model Testing, Validation & Benchmarking

Rank	Tool / Framework	Description
1	<b>Scikit-learn Metrics</b>	Standard metrics (accuracy, F1, RMSE) for ML evaluation.
2	<b>HuggingFace Evaluate</b>	Unified evaluation library for NLP, CV, audio.
3	<b>MLflow Evaluation</b>	Centralized model performance tracking.
4	<b>TensorBoard</b>	Visualize training curves, metrics, embeddings.
5	<b>Weights &amp; Biases (W&amp;B)</b>	Experiment tracking and evaluation dashboards.
6	<b>DeepEval</b>	LLM & GenAI evaluation (semantic correctness, hallucinations).
7	<b>OpenAI Evals</b>	Evaluate LLM outputs against custom test cases.
8	<b>EleutherAI LM Evaluation Harness</b>	Standard benchmarking suite for LLMs.
9	<b>Robustness Gym</b>	Stress testing ML models under different conditions.
10	<b>Great Expectations</b>	Data quality & validation framework.

---

## 11 — OPTIMIZATION – Speed, Memory, Quantization, Fine-Tuning

Rank	Tool / Framework	Description
1	<b>ONNX Runtime</b>	Optimized inference runtime for CPU/GPU.
2	<b>TensorRT</b>	NVIDIA GPU high-speed model optimization.
3	<b>OpenVINO</b>	Intel hardware optimization & quantization.
4	<b>Bitsandbytes</b>	4-bit & 8-bit quantization for LLMs.
5	<b>PEFT (LoRA / QLoRA)</b>	Parameter-efficient fine-tuning.
6	<b>TorchScript</b>	JIT compilation for PyTorch models.
7	<b>XLA (Accelerated Linear Algebra)</b>	Optimization backend for TensorFlow/JAX.
8	<b>DeepSpeed</b>	Large-model optimization & training acceleration.

9	<b>FlashAttention</b>	High-performance attention kernel for LLMs.
10	<b>Graph Optimization (TF Lite)</b>	Mobile/edge optimization tools.

---

## 12 — DEPLOYMENT – Serving, Hosting, Productionization

Rank	Tool / Framework	Description
1	<b>FastAPI</b>	High-performance API serving for ML models.
2	<b>Docker</b>	Containerization for scalable deployments.
3	<b>Kubernetes</b>	Container orchestration for production ML.
4	<b>TensorFlow Serving</b>	High-throughput TF model serving.
5	<b>TorchServe</b>	PyTorch native model server.
6	<b>HuggingFace TGI</b>	Production inference server for LLMs.
7	<b>Seldon Core</b>	Kubernetes-native model serving platform.
8	<b>BentoML</b>	Unified serving framework for all ML models.
9	<b>NVIDIA Triton Inference Server</b>	High-performance multi-framework serving.
10	<b>VLLM Server</b>	Fastest deployment server for large language models.

---

## 13 — MLOps – End-to-End Machine Learning Operations

Rank	Tool / Framework	Description
1	<b>MLflow</b>	Experiment tracking, models, and deployments.
2	<b>Kubeflow</b>	Kubernetes-native MLOps pipelines.
3	<b>Airflow</b>	ML workflow orchestration and scheduling.
4	<b>DVC (Data Version Control)</b>	Dataset & model versioning.
5	<b>Weights &amp; Biases</b>	Experiment tracking, metrics, MLOps dashboards.
6	<b>ZenML</b>	Production-grade MLOps framework.
7	<b>Prefect</b>	Modern workflow orchestration for ML pipelines.
8	<b>Argo Workflows</b>	Kubernetes-native pipeline automation.

9	<b>Neptune.ai</b>	Model registry & monitoring.
10	<b>AWS SageMaker</b>	Full ML lifecycle platform (training → deploy).

---

## 14 — HARDWARE ACCELERATION – GPU/TPU/AI Chips

Rank	Hardware / Framework	Description
1	<b>CUDA</b>	NVIDIA GPU computing backbone.
2	<b>cuDNN</b>	NVIDIA deep neural network acceleration library.
3	<b>TensorRT</b>	GPU inference accelerator for NVIDIA devices.
4	<b>OpenVINO Toolkit</b>	Intel hardware acceleration & optimization.
5	<b>ROCm (AMD)</b>	Open compute platform for AMD GPUs.
6	<b>Google TPU / XLA</b>	High-performance tensor processing for training.
7	<b>Apple CoreML</b>	On-device ML acceleration for iOS/Mac.
8	<b>Qualcomm AI Engine</b>	On-device AI acceleration for mobile chips.
9	<b>NPU (AI Accelerators)</b>	Dedicated neural processors on modern devices.
10	<b>FPGA / VHDL Accelerators</b>	Custom hardware acceleration for enterprise AI.